

# Cs559/659 Lecture 8: Classification Learning

**Readings:** Bishop: Chapter 4.1-2.

# Binary classification

- **Two classes**  $Y = \{0,1\}$
- Our goal is to learn to classify correctly two types of examples
  - Class 0 – labeled as 0,
  - Class 1 – labeled as 1
- We would like to learn  $f : X \rightarrow \{0,1\}$
- **Zero-one error (loss) function**

$$Error_1(\mathbf{x}_i, y_i) = \begin{cases} 1 & f(\mathbf{x}_i, \mathbf{w}) \neq y_i \\ 0 & f(\mathbf{x}_i, \mathbf{w}) = y_i \end{cases}$$

- Error we would like to minimize:  $E_{(x,y)}(Error_1(\mathbf{x}, y))$
- **First step:** we need to devise a model of the function

## **Evaluation of classifiers**

## Evaluation

For any data set we use to test the classification model on we can build a **confusion matrix**:

- Counts of examples with:
- class label  $\omega_j$  that are classified with a label  $\alpha_i$

		target	
		$\omega = 1$	$\omega = 0$
predict	$\alpha = 1$	140	17
	$\alpha = 0$	20	54

# Evaluation

For any data set we use to test the model we can build a **confusion matrix**:

		target	
		$\omega = 1$	$\omega = 0$
predict	$\alpha = 1$	140	17
	$\alpha = 0$	20	54

agreement

Error: ?

# Evaluation

For any data set we use to test the model we can build a confusion matrix:

		target	
		$\omega = 1$	$\omega = 0$
predict	$\alpha = 1$	140	17
	$\alpha = 0$	20	54

agreement

**Error:**  $= 37/231$

**Accuracy**  $= 1 - \text{Error} = 194/231$

# Evaluation for binary classification

Entries in the confusion matrix for binary classification have names:

		target	
		$\omega = 1$	$\omega = 0$
predict	$\alpha = 1$	<i>TP</i>	<i>FP</i>
	$\alpha = 0$	<i>FN</i>	<i>TN</i>

*TP: True positive (hit)*

*FP: False positive (false alarm)*

*TN: True negative (correct rejection)*

*FN: False negative (a miss)*

## Additional statistics

- Sensitivity (recall)

$$SENS = \frac{TP}{TP + FN}$$

- Specificity

$$SPEC = \frac{TN}{TN + FP}$$

- Positive predictive value (precision)

$$PPT = \frac{TP}{TP + FP}$$

- Negative predictive value

$$NPV = \frac{TN}{TN + FN}$$



## Binary classification: additional statistics


- Confusion matrix

		target		
		1	0	
predict	1	140	10	$PPV=140/150$
	0	20	180	$NPV=180/200$
		$SENS=140/160$ $SPEC=180/190$		

### Row and column quantities:

- Sensitivity (SENS)
- Specificity (SPEC)
- Positive predictive value (PPV)
- Negative predictive value (NPV)

## Binary decisions: Receiver Operating Curves



- **Probabilities:**

- *SENS*

$$p(x > x^* \mid \mathbf{X} \in \omega_2)$$

- *SPEC*

$$p(x < x^* \mid \mathbf{X} \in \omega_1)$$


# Receiver Operating Characteristic (ROC)

- ROC curve plots :

$$SN = p(x > x^* | \mathbf{x} \in \omega_1)$$

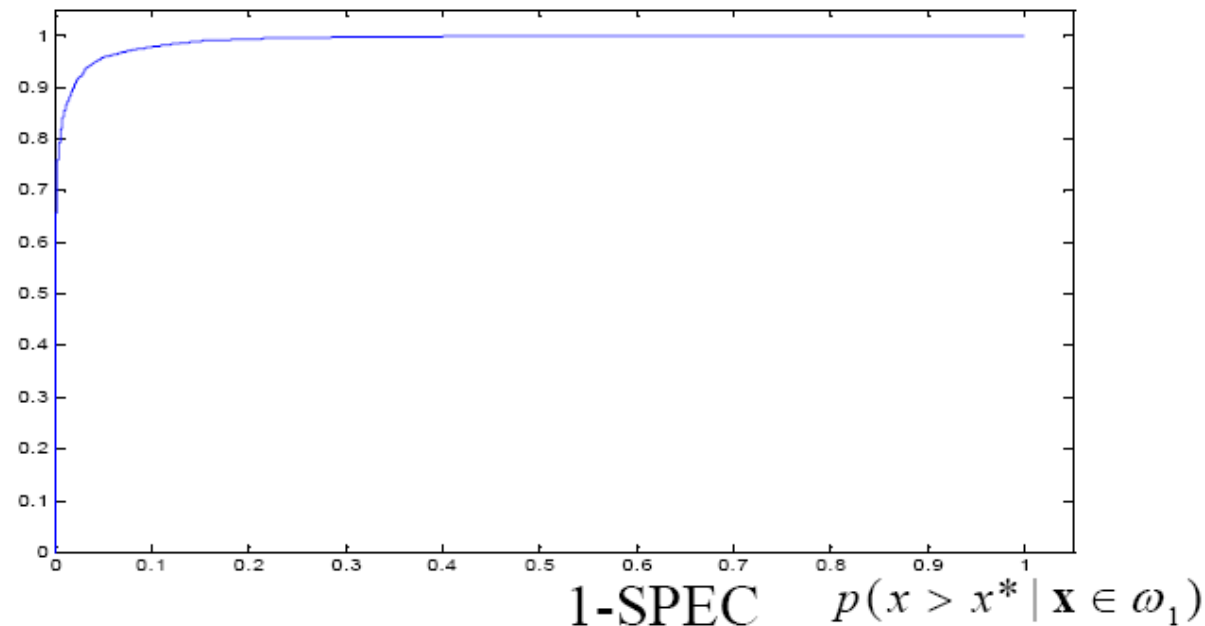
$$1-SP = p(x > x^* | \mathbf{x} \in \omega_2)$$

for different  $x^*$




SENS

$$p(x > x^* | \mathbf{x} \in \omega_2)$$



# ROC curve



# Receiver operating characteristic


- **ROC**
  - shows the discriminability between the two classes under different decision biases
- **Decision bias**
  - can be changed using different loss function

**Back to classification models**

## Discriminant functions


- One way to represent a **classifier** is by using
  - **Discriminant functions**
- **Works for binary and multi-way classification**
- **Idea:**
  - For every class  $i = 0, 1, \dots, k$  define a function  $g_i(\mathbf{x})$  mapping  $X \rightarrow \mathcal{R}$
  - When the decision on input  $\mathbf{x}$  should be made choose the class with the highest value of  $g_i(\mathbf{x})$
- So what happens with the input space? Assume a binary case.

# Discriminant functions






# Discriminant functions

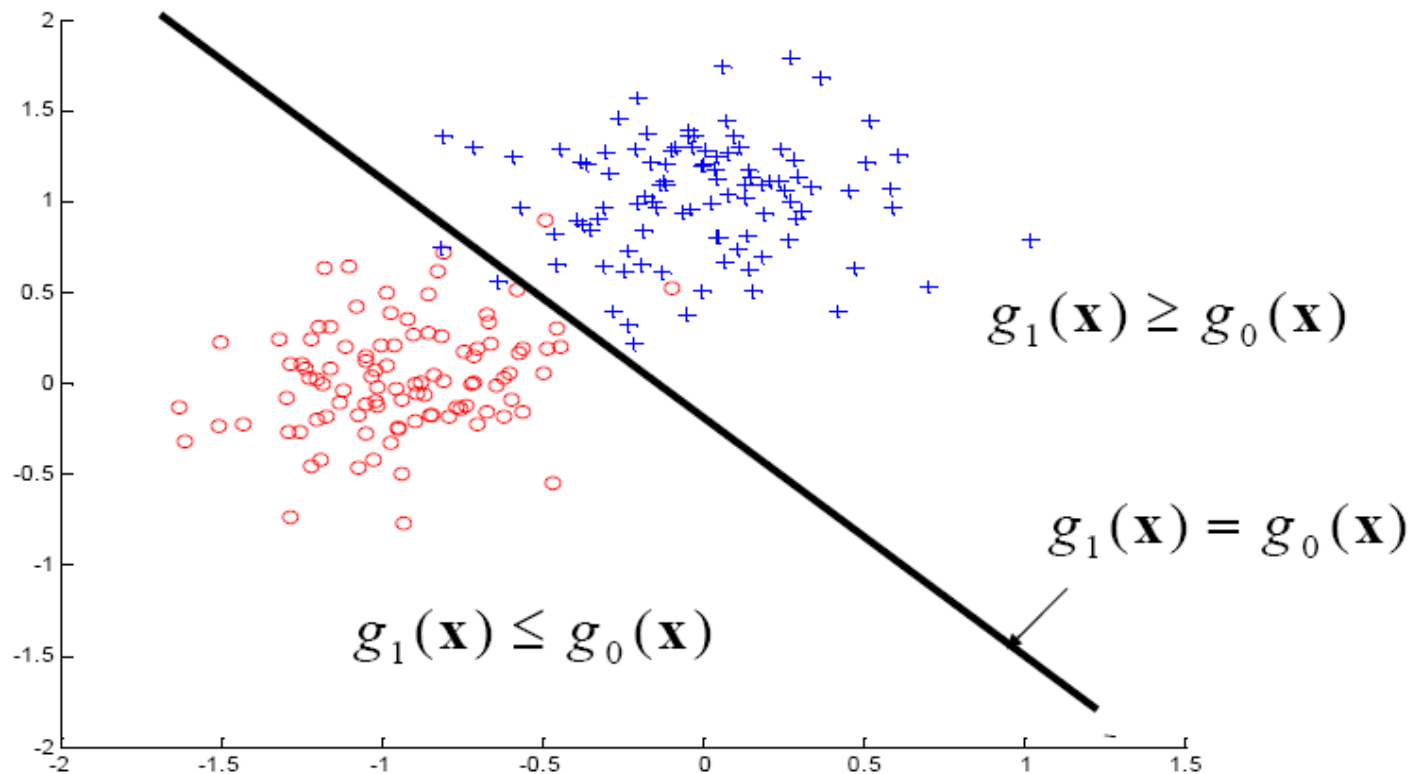


# Discriminant functions




# Discriminant functions

- Define **decision boundary**



# Quadratic decision boundary




# Logistic regression model

- Defines a linear decision boundary
- Discriminant functions:

$$g_1(\mathbf{x}) = g(\mathbf{w}^T \mathbf{x}) \qquad g_0(\mathbf{x}) = 1 - g(\mathbf{w}^T \mathbf{x})$$

- **where**  $g(z) = 1/(1 + e^{-z})$  - is a logistic function

$$f(\mathbf{x}, \mathbf{w}) = g_1(\mathbf{w}^T \mathbf{x}) = g(\mathbf{w}^T \mathbf{x})$$




# Logistic function

function

$$g(z) = \frac{1}{(1 + e^{-z})}$$

- Is also referred to as a **sigmoid function**
- Replaces the threshold function with smooth switching
- takes a real number and outputs the number in the interval  $[0,1]$




# Logistic regression model

- **Discriminant functions:**

$$g_1(\mathbf{x}) = g(\mathbf{w}^T \mathbf{x}) \qquad g_0(\mathbf{x}) = 1 - g(\mathbf{w}^T \mathbf{x})$$

- **Values of discriminant functions vary in  $[0,1]$** 
  - **Probabilistic interpretation**


$$f(\mathbf{x}, \mathbf{w}) = p(y = 1 \mid \mathbf{w}, \mathbf{x}) = g_1(\mathbf{x}) = g(\mathbf{w}^T \mathbf{x})$$



## Logistic regression model. Decision boundary

- LR defines a linear decision boundary

**Example:** 2 classes (blue and red points)





# Generative approach to classification

**Idea:**

1. **Represent and learn the distribution**  $p(\mathbf{x}, y)$
2. **Use it to define probabilistic discriminant functions**

**E.g.**  $g_0(\mathbf{x}) = p(y = 0 | \mathbf{x})$      $g_1(\mathbf{x}) = p(y = 1 | \mathbf{x})$

**Typical model**     $p(\mathbf{x}, y) = p(\mathbf{x} | y) p(y)$

- $p(\mathbf{x} | y) =$  **Class-conditional distributions (densities)**


binary classification: two class-conditional distributions

$$p(\mathbf{x} | y = 0) \qquad p(\mathbf{x} | y = 1)$$

- $p(y) =$  **Priors on classes** - probability of class  $y$

binary classification: Bernoulli distribution

$$p(y = 0) + p(y = 1) = 1$$



# Quadratic discriminant analysis (QDA)

**Model:**

- **Class-conditional distributions**
  - **multivariate normal distributions**

$$\mathbf{x} \sim N(\boldsymbol{\mu}_0, \boldsymbol{\Sigma}_0) \quad \text{for } y = 0$$


$$\mathbf{x} \sim N(\boldsymbol{\mu}_1, \boldsymbol{\Sigma}_1) \quad \text{for } y = 1$$

Multivariate normal  $\mathbf{x} \sim N(\boldsymbol{\mu}, \boldsymbol{\Sigma})$


$$p(\mathbf{x} \mid \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{(2\pi)^{d/2} |\boldsymbol{\Sigma}|^{1/2}} \exp \left[ -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right]$$

- **Priors on classes (class 0,1)**  $y \sim \text{Bernoulli}$ 
  - **Bernoulli distribution**


$$p(y, \theta) = \theta^y (1 - \theta)^{1-y} \quad y \in \{0, 1\}$$




# QDA



## 2 Gaussian class-conditional densities




# QDA: Quadratic decision boundary




# Linear discriminant analysis (LDA)


- When covariances are the same  $\mathbf{x} \sim N(\boldsymbol{\mu}_0, \boldsymbol{\Sigma}), y = 0$   
 $\mathbf{x} \sim N(\boldsymbol{\mu}_1, \boldsymbol{\Sigma}), y = 1$



# LDA: Linear decision boundary



# LDA: linear decision boundary





# Logistic regression vs LDA

- **Two models with linear decision boundaries:**
  - **Logistic regression**
  - **Generative model with 2 Gaussians with the same covariance matrices**

$$x \sim N(\mu_0, \Sigma) \quad \text{for} \quad y = 0$$

$$x \sim N(\mu_1, \Sigma) \quad \text{for} \quad y = 1$$

- **Two models are related !!!**
  - When we have **2 Gaussians with the same covariance matrix** the probability of  $y$  given  $\mathbf{x}$  has the form of a logistic regression model !!!

$$p(y = 1 \mid \mathbf{x}, \mu_0, \mu_1, \Sigma) = g(\mathbf{w}^T \mathbf{x})$$

# When is the logistic regression model correct?

- **Members of the exponential family can be often more naturally described as**

$$f(\mathbf{x} \mid \boldsymbol{\theta}, \boldsymbol{\varphi}) = h(x, \boldsymbol{\varphi}) \exp \left\{ \frac{\boldsymbol{\theta}^T \mathbf{x} - A(\boldsymbol{\theta})}{a(\boldsymbol{\varphi})} \right\}$$


$\boldsymbol{\theta}$  - A location parameter       $\boldsymbol{\varphi}$  - A scale parameter

- **Claim:** A logistic regression is a correct model when class conditional densities are from the same distribution in the exponential family and have **the same scale factor**  $\boldsymbol{\varphi}$
- **Very powerful result !!!!**
  - **We can represent posteriors of many distributions with the same small network**

# Linear units

## Linear regression

$$f(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$$




### Gradient update:

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha \sum_{i=1}^n (y_i - f(\mathbf{x}_i)) \mathbf{x}_i$$

Online:  $\mathbf{w} \leftarrow \mathbf{w} + \alpha (y - f(\mathbf{x})) \mathbf{x}$

## Logistic regression

$$f(\mathbf{x}) = p(y = 1 | \mathbf{x}, \mathbf{w}) = g(\mathbf{w}^T \mathbf{x})$$



### Gradient update:

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha \sum_{i=1}^n (y_i - f(\mathbf{x}_i)) \mathbf{x}_i$$

Online:  $\mathbf{w} \leftarrow \mathbf{w} + \alpha (y - f(\mathbf{x})) \mathbf{x}$

**The same**



# Gradient-based learning

- The **same simple gradient update rule** derived for both the linear and logistic regression models
- Where the magic comes from?
- Under the **log-likelihood** measure the function models and the models for the output selection fit together:

- **Linear model + Gaussian noise**



$$y = \mathbf{w}^T \mathbf{x} + \varepsilon$$

$$\varepsilon \sim N(0, \sigma^2)$$

- **Logistic + Bernoulli**

$$y \sim \text{Bern}(\theta)$$

$$\theta = p(y = 1 | \mathbf{x}) = g(\mathbf{w}^T \mathbf{x})$$



# Generalized linear models (GLIM)

## Assumptions:

- The conditional mean (expectation) is:  $\mu = f(\mathbf{w}^T \mathbf{x})$ 
  - $f(\cdot)$  is a **response (or a link) function**
- Output  $y$  is characterized by an exponential family distribution with mean  $\mu = f(\mathbf{w}^T \mathbf{x})$

## Examples:

### – Linear model + Gaussian noise



$$y = \mathbf{w}^T \mathbf{x} + \varepsilon \quad \varepsilon \sim N(0, \sigma^2)$$

$$y \sim N(\mathbf{w}^T \mathbf{x}, \sigma^2)$$

### – Logistic + Bernoulli

$$y \sim \text{Bern}(\theta) \sim \text{Bern}(g(\mathbf{w}^T \mathbf{x}))$$

$$\theta = g(\mathbf{w}^T \mathbf{x}) = \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}}}$$



# Generalized linear models (GLMs)

- **A canonical response functions  $f(.)$  :**
  - **encoded in the distribution**

$$p(\mathbf{x} \mid \boldsymbol{\theta}, \boldsymbol{\varphi}) = h(x, \boldsymbol{\varphi}) \exp \left\{ \frac{\boldsymbol{\theta}^T \mathbf{x} - A(\boldsymbol{\theta})}{a(\boldsymbol{\varphi})} \right\}$$

- **Leads to a simple gradient form**
- **Example: Bernoulli distribution**

$$p(x \mid \mu) = \mu^x (1 - \mu)^{1-x} = \exp \left\{ \log \left( \frac{\mu}{1 - \mu} \right) x + \log(1 - \mu) \right\}$$

$$\theta = \log \left( \frac{\mu}{1 - \mu} \right) \qquad \mu = \frac{1}{1 + e^{-\theta}}$$

- **Logistic function matches the Bernoulli**

# Non-linear extension of logistic regression

- use **feature (basis) functions** to model **nonlinearities**
  - the same trick as used for the linear regression

## Linear regression

$$f(\mathbf{x}) = w_0 + \sum_{j=1}^m w_j \phi_j(\mathbf{x})$$

$\phi_j(\mathbf{x})$  - an arbitrary function of  $\mathbf{x}$

## Logistic regression

$$f(\mathbf{x}) = g\left(w_0 + \sum_{j=1}^m w_j \phi_j(\mathbf{x})\right)$$

