

CS559 Lecture 12: Gaussian Process

Reading: Chapter 6, Bishop book

Gaussian Process

- Extension of kernels to probabilistic discriminant models
- Define a prior probability distribution over the functions directly.
- Only need to consider function values at training and test data sets – work in finite space.
- Kriging, ARMA, Kalman filters, RBFs are all Gaussian Processes.
- <http://www.gaussianprocess.org/>

Gaussian Stochastic Process

Consider linear model

$$y(x) = w^T \phi(x)$$

with prior over w

$$p(w) = \mathcal{N}(w|0, \alpha^{-1}I)$$

Now, if \mathbf{y} is vector of samples from a stochastic process, then it is a Gaussian stochastic process with

$$E[\mathbf{y}] = \Phi E[w] = 0$$


$$E[\mathbf{y}] = \Phi E[ww^T] \Phi^T = \frac{1}{\alpha} \Phi \Phi^T = K$$

where K is the Gram matrix

$$K_{nm} = k(x_n, x_m) = \frac{1}{\alpha} \phi(x_n)^T \phi(x_m)$$

Gaussian process

- Gaussian process is a probability distribution over function $y(x)$ such that the values of y 's at points x_1, \dots, x_N have a joint Gaussian distribution.
- 2D -> Gaussian random field.



$$k_1(x_n, x_m) = \exp(-||x_n - x_m||^2 / 2\sigma^2)$$

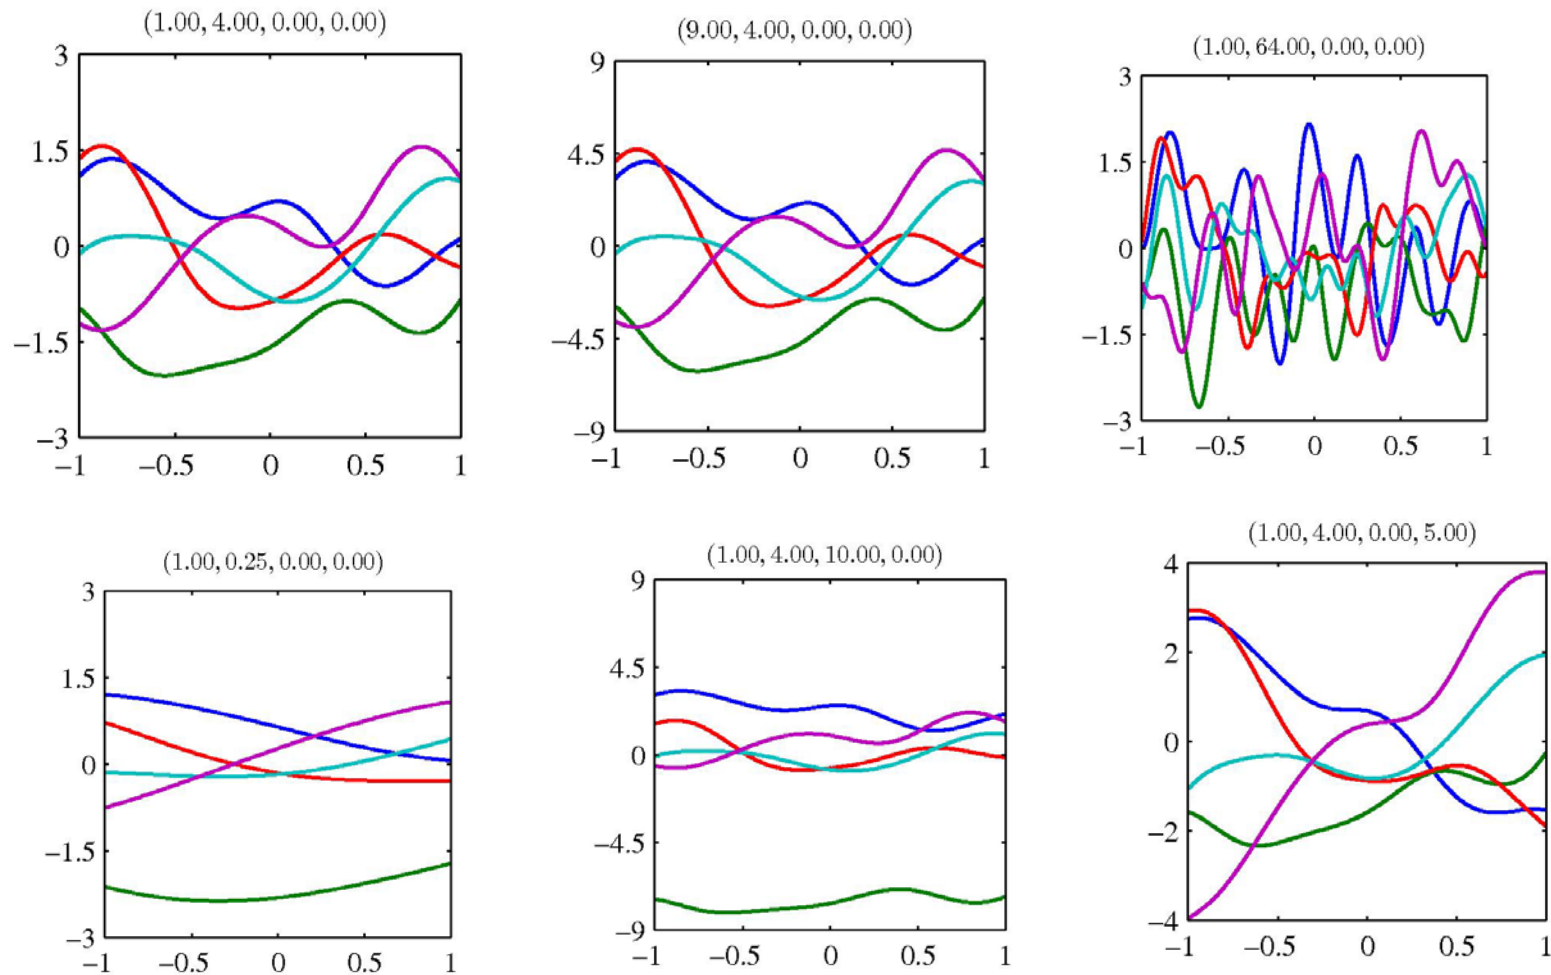
$$k_2(x_n, x_m) = \exp(-\theta |x_n - x_m|)$$

Gaussian stochastic processes are completely defined by the second order statistics!

Gaussian Process

One popular choice of kernel in this case is

$$k(x_n, x_m) = \theta_0 \exp \left(-\frac{\theta_1}{2} \|x_n - x_m\|^2 \right) + \theta_2 + \theta_3 x_n^T x_m$$



Gaussian Process for Regression

$$t_n = y_n + \epsilon_n, \quad \epsilon_n \sim \mathcal{N}(0, \beta^{-1})$$

Since y_n and ϵ_n are independent,

$$p(\mathbf{t}|\mathbf{y}) = \mathcal{N}(\mathbf{t}|\mathbf{y}, \beta^{-1}\mathbf{I})$$

and we know


$$p(\mathbf{y}) = \mathcal{N}(\mathbf{y}|0, K)$$

Therefore

$$p(\mathbf{t}) = \int p(\mathbf{t}|\mathbf{y})p(\mathbf{y})d\mathbf{y} = \mathcal{N}(\mathbf{t}|0, \mathbf{C})$$

where

$$C(x_n, x_m) = k(x_n, x_m) + \beta^{-1}\delta_{nm}$$



Gaussian Process for Regression

For regression, however, we need $p(t_{n+1}|\mathbf{t})$. We start with

$$p(\mathbf{t}_{n+1}) = \mathcal{N}(\mathbf{t}_{n+1}|0, \mathbf{C}_{n+1})$$



where \mathbf{C}_{n+1} is a $(N+1) \times (N+1)$ covariance matrix.

Since the underlying process is a Gaussian stochastic process, you can marginalize the distribution by partitioning the covariance (2.81,2.82), thus

$$\begin{aligned} E[t_{n+1}] &= m(x_{n+1}) = \mathbf{k}^T \mathbf{C}_n^{-1} \mathbf{t} \\ E[t_{n+1}t_{n+1}] &= \sigma^2(x_{n+1}) = c - \mathbf{k}^T \mathbf{C}_n^{-1} \mathbf{k} \end{aligned}$$

$m(x_{n+1})$ can also be written as

$$m(x_{n+1}) = \sum_{n=1}^N a_n k(x_n, x_{n+1})$$



If the kernel function is chosen as a specific finite set of basis functions, it will lead to the Linear Regression case we talked about before.

Parametric space + linear regression \leftrightarrow functional space + Gaussian Process

Multilayer neural networks


**Or another way of modeling nonlinearities
for regression and classification problems**

Reading: Chapter 5, Bishop book.

Linear units

Linear regression

$$f(\mathbf{x}) = w_0 + \sum_{j=1}^d w_j x_j$$




On-line gradient update:

$$\begin{aligned} w_0 &\leftarrow w_0 + \alpha(y - f(\mathbf{x})) \\ &\vdots \\ w_j &\leftarrow w_j + \alpha(y - f(\mathbf{x}))x_j \end{aligned}$$

Logistic regression

$$f(\mathbf{x}) = p(y=1 | \mathbf{x}, \mathbf{w}) = g(w_0 + \sum_{j=1}^d w_j x_j)$$



On-line gradient update:

$$\begin{aligned} w_0 &\leftarrow w_0 + \alpha(y - f(\mathbf{x})) \\ &\vdots \\ w_j &\leftarrow w_j + \alpha(y - f(\mathbf{x}))x_j \end{aligned}$$


The same



Limitations of basic linear units

Linear regression


$$f(\mathbf{x}) = w_0 + \sum_{j=1}^d w_j x_j$$



Function linear in inputs !!

Logistic regression

$$f(\mathbf{x}) = p(y=1 | \mathbf{x}, \mathbf{w}) = g(w_0 + \sum_{j=1}^d w_j x_j)$$




Linear decision boundary!!

Regression with the quadratic model.


Limitation: linear hyper-plane only

- a non-linear surface can be better




Linear decision boundary

- logistic regression model is not optimal, but not that bad




When logistic regression fails?

- Example in which the logistic regression model fails



Limitations of linear units.

- Logistic regression does not work for **parity functions**
 - no linear decision boundary exists



Solution: a model of a non-linear decision boundary

Extensions of simple linear units

- use **feature (basis) functions** to model **nonlinearities**


Linear regression

$$f(\mathbf{x}) = w_0 + \sum_{j=1}^m w_j \phi_j(\mathbf{x})$$

Logistic regression

$$f(\mathbf{x}) = g\left(w_0 + \sum_{j=1}^m w_j \phi_j(\mathbf{x})\right)$$

$\phi_j(\mathbf{x})$ - an arbitrary function of \mathbf{x}



Learning with extended linear units


Feature (basis) functions model **nonlinearities**

Linear regression

$$f(\mathbf{x}) = w_0 + \sum_{j=1}^m w_j \phi_j(\mathbf{x})$$

Logistic regression

$$f(\mathbf{x}) = g\left(w_0 + \sum_{j=1}^m w_j \phi_j(\mathbf{x})\right)$$




Important property:

- The same problem as learning of the weights for linear units, the input has changed— but the weights are linear in the new input

Problem: too many weights to learn

Multi-layered neural networks


- An alternative way to introduce **nonlinearities to regression/classification models**
- **Key idea: Cascade several simple neural models with logistic units.** Much like neuron connections.





- ▶ Perceptrons \equiv nerve cells or neurons
- ▶ Network of perceptrons \equiv connections between neurons

Biological Neural Network

The human brain has approximately 100 billion nerve cells, called neurons, each connected to thousands of others. Senses and thoughts trigger electrical impulses that quickly travel through the neural network. When a neuron receives information, it can send the message on or stop it from travelling forward.



Neurons vs. Perceptrons




But, there are real differences, firing rate vs. potential, etc

Multilayer neural network

Also called a **multilayer perceptron (MLP)**

Cascades multiple logistic regression units

Example: (2 layer) classifier with non-linear decision boundaries




Input layer

Hidden layer

Output layer


Multilayer neural network

- Models **non-linearity through logistic regression units**
- Can be applied to both **regression and binary classification problems**



Multilayer neural network

- **Non-linearities are modeled using multiple hidden logistic regression units (organized in layers)**
- The output layer determines whether it is a **regression** or a **binary classification problem**



Expressive Power of NNs

- ▶ Can NNs express any function $\mathcal{F} : \mathbb{R}^d \rightarrow \mathbb{R}$?
- ▶ Yes, approximately!

Kolmogorov's Superposition Theorem (1957)

For all $n \geq 2$ and for any continuous real function \mathcal{F} of n variables in the domain $[0, 1]$, $\mathcal{F} : [0, 1]^n \mapsto \mathbb{R}$, there exists $(2n + 1)n$ continuous, monotone increasing univariate functions on $[0, 1]$, by which \mathcal{F} can be reconstructed as:

$$\mathcal{F}(x_1, \dots, x_n) = \sum_{q=1}^{2n+1} \Phi_q \left(\sum_{p=1}^n \Psi_{pq}(x_p) \right)$$



Universal Approximation Theorem (Hornik 1991)

Any continuous, nonconstant, bounded and monotonically increasing function, \mathcal{F} , can be approximated arbitrarily closely, $\epsilon > 0$, by a feed-forward network with one semilinear hidden units using a threshold function and one linear output unit, formally,

$$\max \left| \mathcal{F}(x_1, \dots, x_n) - \sum_{k=1}^K w_k^{(2)} \Phi \left(\sum_{j=1}^n w_{kj}^{(1)} x_j - w_{0j}^{(1)} \right) \right| < \epsilon$$

- ▶ Semilinear units: $g(L(x) - b)$, where $L(x)$ is linear in x and g (hard limiter) is a monotone real function with the limits

$$\lim_{x \rightarrow -\infty} g(x) = 0 \quad \text{and} \quad \lim_{x \rightarrow \infty} g(x) = 1$$

- ▶ This approximation is lenient about activation functions

Learning in NNs


- ▶ Okay, so NNs have tremendous expressive power
- ▶ But, is there a way to learn the function for a given problem?
- ▶ Yes, the back propagation algorithm!

Learning with MLP

- How to learn the parameters of the neural network?
- **Gradient descent algorithm**
 - Weight updates based on the error: $J(D, \mathbf{w})$

$$\mathbf{w} \leftarrow \mathbf{w} - \alpha \nabla_{\mathbf{w}} J(D, \mathbf{w})$$

- We need to **compute gradients for weights in all units**
- **Can be computed in one backward sweep through the net !!!**




- The process is called **back-propagation**

Backpropagation

(k-1)-th level

k-th level

(k+1)-th level



$x_i(k)$ - output of the unit i on level k

$z_i(k)$ - input to the sigmoid function on level k

$w_{i,j}(k)$ - weight between units j and i on levels (k-1) and k

$$z_i(k) = w_{i,0}(k) + \sum_j w_{i,j}(k)x_j(k-1)$$

$$x_i(k) = g(z_i(k))$$

Backpropagation

Update weight $w_{i,j}(k)$ using a data point $D = \{< \mathbf{x}, y >\}$

$$w_{i,j}(k) \leftarrow w_{i,j}(k) - \alpha \frac{\partial}{\partial w_{i,j}(k)} J(D, \mathbf{w})$$

Let $\delta_i(k) = \frac{\partial}{\partial z_i(k)} J(D, \mathbf{w})$

Then: $\frac{\partial}{\partial w_{i,j}(k)} J(D, \mathbf{w}) = \frac{\partial J(D, \mathbf{w})}{\partial z_i(k)} \frac{\partial z_i(k)}{\partial w_{i,j}(k)} = \delta_i(k) x_j(k-1)$

S.t. $\delta_i(k)$ is computed from $x_i(k)$ and the next layer $\delta_l(k+1)$

$$\delta_i(k) = \left[\sum_l \delta_l(k+1) w_{l,i}(k+1) \right] x_i(k) (1 - x_i(k))$$

Last unit (is the same as for the regular linear units):

$$\delta_i(K) = - \sum_{u=1}^n (y_u - f(\mathbf{x}_u, \mathbf{w}))$$

It is the same for the classification with the log-likelihood measure of fit and linear regression with least-squares error!!!

Learning with MLP

- **Gradient descent algorithm**

- Weight update:

$$w_{i,j}(k) \leftarrow w_{i,j}(k) - \alpha \frac{\partial}{\partial w_{i,j}(k)} J(D, \mathbf{w})$$

$$\frac{\partial}{\partial w_{i,j}(k)} J(D, \mathbf{w}) = \frac{\partial J(D, \mathbf{w})}{\partial z_i(k)} \frac{\partial z_i(k)}{\partial w_{i,j}(k)} = \delta_i(k) x_j(k-1)$$

$$w_{i,j}(k) \leftarrow w_{i,j}(k) - \alpha \delta_i(k) x_j(k-1)$$

$x_j(k-1)$ - j-th output of the (k-1) layer

$\delta_i(k)$ - derivative computed via back-propagation

α - a learning rate

Learning with MLP

- **Online gradient descent algorithm**

- Weight update:

$$w_{i,j}(k) \leftarrow w_{i,j}(k) - \alpha \frac{\partial}{\partial w_{i,j}(k)} J_{\text{online}}(D_u, \mathbf{w})$$

$$\frac{\partial}{\partial w_{i,j}(k)} J_{\text{online}}(D_u, \mathbf{w}) = \frac{\partial J_{\text{online}}(D_u, \mathbf{w})}{\partial z_i(k)} \frac{\partial z_i(k)}{\partial w_{i,j}(k)} = \delta_i(k) x_j(k-1)$$

$$w_{i,j}(k) \leftarrow w_{i,j}(k) - \alpha \delta_i(k) x_j(k-1)$$

$x_j(k-1)$ - j-th output of the (k-1) layer

$\delta_i(k)$ - derivative computed via backpropagation

α - a learning rate

Online gradient descent algorithm for MLP

Online-gradient-descent (D , *number of iterations*)

Initialize all weights $w_{i,j}(k)$

for $i=1:1:$ *number of iterations*

do **select** a data point $D_u = \langle \mathbf{x}, y \rangle$ from D

set learning rate α

compute outputs $x_j(k)$ for each unit

compute derivatives $\delta_i(k)$ via **backpropagation**

update all weights (in parallel)


$$w_{i,j}(k) \leftarrow w_{i,j}(k) - \alpha \delta_i(k) x_j(k-1)$$

end for


return weights \mathbf{w}

Xor Example.

- linear decision boundary does not exist




Xor example. Linear unit




Xor example.

Neural network with 2 hidden units



Xor example.

Neural network with 10 hidden units



MLP in practice

- **Optical character recognition** – digits 20x20
 - Automatic sorting of mails
 - 5 layer network with multiple output functions

